

Wizbit, a distributed revisioning file system

**Rob Taylor and Mark Doffman
Codethink Ltd**



Contents

- **Background**
 - File system problems & solutions
- **Goals**
 - What we intend to solve
- **Design**
 - Git internals & implementation
 - Wizbit design
- **Blinding future**
 - How Wizbit may be (mis)used in the future

Background



What's the problem?

- **We all now have lots of devices**
 - Computers
 - Web services
 - Phones
 - Tivos
 - Tablets
 -

What's the problem??

- **All of these have their own file/data stores**
 - You need to regularly make decisions about what should be on what.
 - You need to transfer data between the devices
 - Things get out of sync, especially when it's something you're working on. Especially if you're working on it with other people.

What's the problem???

- Sync – **FAIL!**
- Collaboration – **FAIL!**
- Backup – **FAIL!** (ish..)

Doing data sharing now

- **Flickr**
- **E-mailing file versions**
- **Shared drives**
- **Google docs**
- **Google file system (Wuala)**
- **SVN**
- **Git / DVCS (Distributed Revision Control Systems)**

Problems

- **Flickr / Google Docs and DVCS are application specific.**
- **E-Mailing revisions about is incredibly frustrating.**
- **Google file system / Shared drives do not allow off-line collaboration.**

Device data synchronisation

- **OpenSync**
- **Conduit**
- **Multiple commercial software packages, allwaysync, GoodSync.**
- **iTunes**

Problems

- **Limited data formats, proprietary or lots of failure cases**
- **Two-way sync is hard.**



Backup

- **Files are not automatically backed up leading to data loss.**
- **Many enterprise solutions.**
- **rsync and dirvish may provide complete solution.**
- **Lack of open-source GUI.**

Where is my data?

- **User experience should be that you have the same file store on all devices**
- **Perhaps different views onto the same file store**
- **Media player gets to see your media files**
- **Desktop sees all**
- **Flickr sees pictures**
- **BUT it should be the same file store!**

So?

- **We need a distributed file system**
- **Must still be fully usable when disconnected**
- **Must cope with conflicting changes**



Conflicts

- **Most existing distributed file systems don't allow conflicts – no disconnected operation possible**
- **Those that do resolve conflicts by a complicated (and error prone) system of rules when the file system syncs (e.g. Coda)**

Conflicts, fixed..

- **We need conflict resolution to involve the user and the application.**
 - Only the applications fully understand their data formats
 - Some merge cases can only be resolved by the user telling the application what is right.
- **When to resolve a conflict should be the choice of the user.**
 - Don't force users to make decisions when they don't need to.
 - Don't get in the way of the workflow

Collaboration

- **If we can solve this, we get free collaboration tools.**



What does the solution look like?

- **Looks a bit like a DVCS (Distributed Revision Control System)**
 - Branching history
 - Merges on request
 - Store synchronization without changing working tree or even necessarily the working branch

What doesn't fit?

- **DVCSs have the luxury of understanding what kind of data they are storing**
 - Text files where unit of merge is a line
- **In a DVCS you merge a whole tree**
 - For a user file system the user has to be able to merge individual files and leave other conflicts unresolved.
- **User interface is not suitable, not even for DVCS ;-)**
 - (ok, maybe with Bazaar..)

Wizbit Design



Attempt 1: Git wrapper design

- **Each file referenced by UUID**
- **Each file stored in separate git repository**
- **XML storage of the directory structure**
- **Python implementation using Popen and git commands**

Attempt 1: **FAIL**

- **Git issues**

- Whole tree merge was unsuitable – Separate repository per file.
- Git repository per file was heavy
- Directory structure is not versioned

- **Implementation issues**

- XML storage of file information was heavy and error prone
- Python and Popen were error prone
- Language bindings and library impossible
- Directory file system is the only application

Attempt 2

- **Move core Git functionality to reusable library**
 - All of Git UI not wanted.
- **Create new directory structure**
 - Uses git object and pack format, but rewrite the storage of references.
- **Build file system on-top of this**
 - File system is a particular application of a versioned collection of files.
- **Git chosen for speed and language choice**
 - 'C' library would make for easier language bindings

Git Overview




Git design

- **Git objects are files referenced by their own hash**
- **They are stored compressed**
- **They may contain data or meta-data**
 - Meta-data in this case is version history.


Git object types

- **Blob**
 - File data objects
- **Tree**
 - Store directory data a list of references to blob or tree objects.
- **Commit**
 - Represent a commit including a commit message, time, references to parent commits, and a single tree object.
- **Tag**
 - Used to implement Git tags, we won't worry about these.

Git tree – Single commit



Git commit tree



Git refs

- **“Entry point” required for commit tree**
- **Hash reference used as this entry point – The most recent commit**



Git design overview

- **All objects referenced by their hash.**
- **A commit object uniquely references and entire directory tree.**
- **Commit tree used for revision history.**
- **Files store branch references**
- **Special file stores merge reference**

Git implementation

- **Written in 'C' with some bash scripts.**
- **Optimized to execute in short live process as a command line application.**
- **An 'evolved' code base that is not obviously modular.**
- **Mostly not reentrant**

Git thread safety

```
/*
 * I'm tired of doing "vsprintf()" etc just to open a
 * file, so here's a "return static buffer with printf"
 * interface for paths.
 *
 * It's obviously not thread-safe. Sue me. But it's quite
 * useful for doing things like
 *
 *   f = open(mkpath("%s/%s.git", base, name), O_RDONLY);
 *
 * which is what it's designed for.
 */
#include "cache.h"

static char bad_path[] = "/bad-path/";

static char *get_pathname(void)
{
    static char pathname_array[4][PATH_MAX];
    static int index;
    return pathname_array[3 & ++index];
}

static char *cleanup_path(char *path)
{
    /* Clean it up */
    if (!memcmp(path, "./", 2)) {
        path += 2;
    }
}
```

- **Thread safety not considered in Git.**
- **Good decision for a command line application.**
- **Concurrency provided by larger git design and file locking.**

Git memory management

- **Memory cleared by OS – process exits!**
- **Frequent memory allocations (Objects) stored in ever-growing cache**
- **Code-base does not consider long-lived processes**

libgit.a suitability

None
(oops!)


So.... libgitcore

- **Reentrant, object-like rewrite of git code functionality (objects, blobs, trees, commits, tags)**
- **Using git code but heavy reworking needed**
- **Usage is one git object-loader and writer per thread.**
- **Git code is horrible..**


Wizbit design overview

- **Files referenced by UUID**
 - Common feature of distributed file systems
 - Ensures that files created on different computers do not have the same reference
- **Commit tree per file**
 - Each file can be merged separately from the rest of the repository

Wizbit data structures



- File tree is a standard format
- Reduced to two blobs
 - /data
 - /meta-data



Wizbit directory structure

- **Commit refs no longer stored in files**
- **Git configuration no longer exists**
- **Object store is the same as git**



Refs in Wizbit - Requirements

- **A single head reference per file**
 - The 'work branch'
- **Multiple other references**
 - Usually unmerged tips from other Wizbit instances
- **Multiple references need to be updated within a single transaction**
 - To provide posix semantics, we'll need to update multiple refs atomically.

Wizbit references - Solution!

- **Use SQLite for the win, avoiding lots of clever file system locking code.**



Difference of two repositories

- **Calculating the differences between repositories is vital for syncing**
- **Checking the head reference of every file is not optimal in the common case of similar repositories**
- **Git / SVN / Everything else do not have this problem as a reference applies to an entire tree**

Difference of two repositories

- **Use a branching tree of UUIDs.**
- **Compare SHA1s at each tree depth**
- **Quickly find out which refs have changed**



Building a file system

- **Options include**
 - Fuse
 - GVFS
- **Choice not important to us – probably do both**
- **Didn't think GVFS had any advantages**

File system structure

- **How do we build a file system on-top of the UUID structure?**
 - Store directories as versioned files in the Wizbit repository
 - Store symlinks as versioned files in the Wizbit repository
 - Store files as versioned files in the Wizbit repository
- **Extremely similar to local file system**
- **UUID could be considered as the inode**

Wizbit design implications

- **Destructively deleting is difficult**
 - Real deleting would require removing commit history.
 - Commit tree modification is nearly impossible on published or shared repositories.
 - For the moment freeing up space will not be supported.
 - Don't use wizbit for your “sensitive data”
- **Checkout of directory revisions is difficult**
 - In the versioned file system directories only point to file uuid rather than file versions, making checking out a directory out at a certain revision non-trivial.

Related Tech

- **Versioned file systems**
 - Andrew FS, CODA, Elephant FS
 - User space, slower but perhaps more flexible
 - Synchronization between repositories built in
- **Log based file systems**
 - User space implementation and commit on close make witzbit unsuitable for the forensics applications of Log based file systems
 - Possibly faster file writing.
 - Possibly better ratio of data to commit data.

Usefulness

- **Not a whole-file system solution**
 - Lack of delete
 - Temporary files to be generally avoided
- **Not a revision control solution**
 - Git / Mercurial / Bazarr much better choices as they are designed to version whole directories + Years of UI development.
- **Intended for personal files, accessed and modified on multiple devices by many people**

Usefulness

- **Intended for creating a single file system across multiple devices**



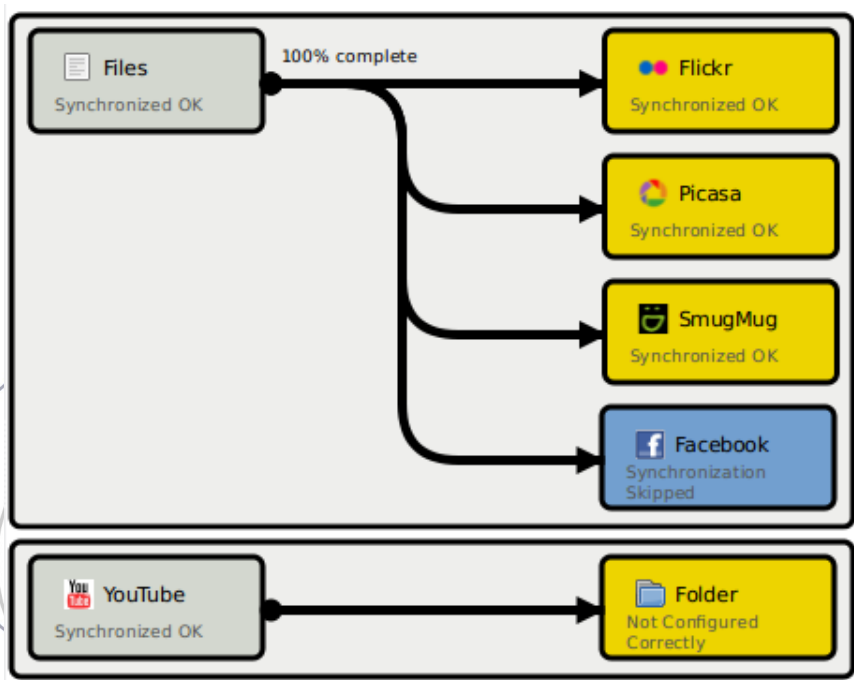
User interface – Local repository

- **File access**
 - Handled by user space file system
- **Committing**
 - When to commit is difficult. Best bet – Commit on close.
- **Version history**
 - Reuse ideas from graphical VCS.
- **Revision checkout**
- **File merging**
 - Application specific and hard. Provide some basics such as “Use most recent file”.

User interface - Syncing

- **Relating devices**
 - Where are we syncing from / to?
- **Sync schedules**
 - When do we sync between devices?
- **Authorization**
 - How do we obtain permission to sync to a device?

Syncing - Conduit



- Aims to integrate into desktop
- Plug and play for your device syncing
- Automatically syncs when something changes
- Already has code to handle authorization against web services

The blinding future!



Meta-data formats

- In the core design, every file has both data and meta-data.
- Not yet decided what form meta-data should take: RDF/XML? Tuples?
- Common meta-data:
 - User visible file name
 - mtime/ctime (atime??!)
 - Document info

Tracker / XESAM

- **Initially can use Tracker to create the meta-data**
- **If Wizbit successfully adopted, applications can create the meta-data themselves and commit along with the data.**

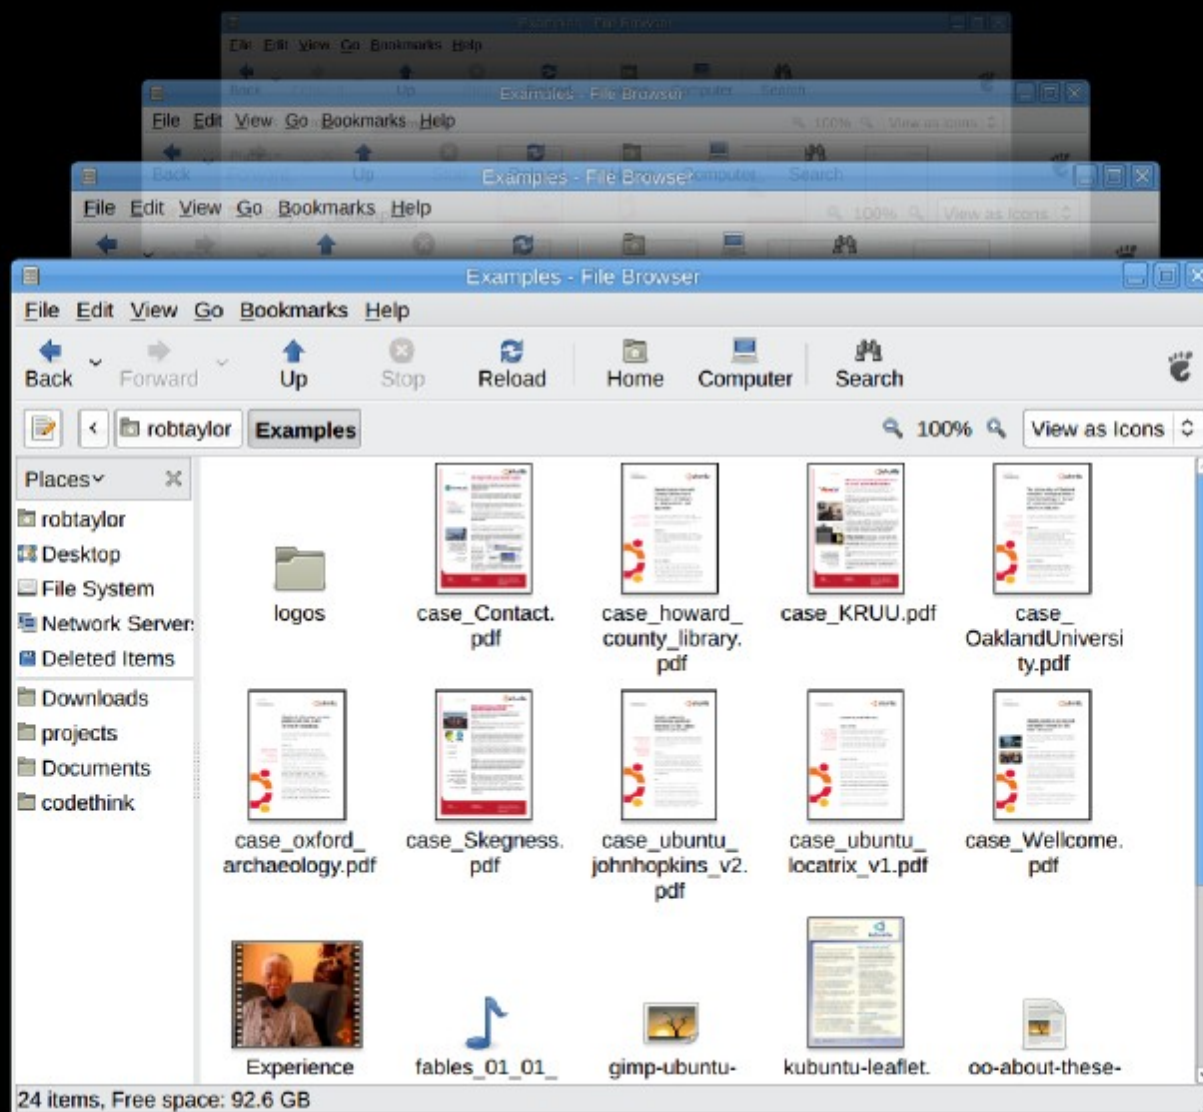

Contextual browsing

- **Given all this and a suitably rich ontology, we can start imagining new ways of reaching our data:**
 - Time line: “The work I was doing last week”
 - Context: “Documents on RDF”
 - Virtual folders with complex stored search terms
 - “Related files”

Metadata

- **Metadata should move with the file system**
 - Shouldn't need to be regenerated for each store
 - What to do about inter-file relations?





Timeline view

- **UI Mockup by Karl Lattimer**
 - Top line shows activity over time – thicker/hotter means more activity
 - Dragging magnifier scrolls time line forwards/backwards
 - Magnifier scale gives detail level
 - Files changing
 - How files changed

More ideas welcome!



Thanks!

- <http://wizbit.org>
 - (not much there yet)
- <http://git.codethink.co.uk>
 - libgitcore
 - wizbit